

Project - 2 Group 16 - Rohit Suseel, Madhumitha Vijayakrishna

Predicting Housing Prices in King County, USA using Regression Analysis

Regression Abstract: This dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015.

Source Link: <https://www.kaggle.com/harlfoxem/housesalesprediction>
(<https://www.kaggle.com/harlfoxem/housesalesprediction>)

MinMaxScaler is used for both the tasks because StandardScaler cannot guarantee balanced feature scales in the presence of outliers.

There are no missing values in the original datasets, values are manually removed.

In [4]:

```
1 import numpy as np #handling numbers
2 import pandas as pd #handling the dataset
3 import matplotlib as mpl
4 import matplotlib.pyplot as plt
5 from sklearn.impute import SimpleImputer # handling missing data
6 from sklearn.preprocessing import LabelEncoder, OneHotEncoder # encoding categorical
7 from sklearn.model_selection import train_test_split # splitting training and testing
8 from sklearn.preprocessing import StandardScaler #feature scaling
9 %matplotlib inline
10 import seaborn as sns
11 import matplotlib.pyplot as plt
12
13 df = pd.read_csv('D:/UTD Fall 2020/AML/Project 1/Project1_Group16/kc_house_data.csv')
14 df.rename(columns={'price': 'SalePrice'}, inplace=True)
15 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 21613 entries, 0 to 21612
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	id	21613 non-null	int64
1	date	21613 non-null	object
2	SalePrice	21613 non-null	float64
3	bedrooms	21613 non-null	int64
4	bathrooms	21613 non-null	float64
5	sqft_living	21613 non-null	int64
6	sqft_lot	21613 non-null	int64
7	floors	21613 non-null	float64
8	waterfront	21613 non-null	int64
9	view	21613 non-null	int64
10	condition	21613 non-null	int64
11	grade	21613 non-null	int64
12	sqft_above	21613 non-null	int64
13	sqft_basement	21613 non-null	int64
14	yr_built	21613 non-null	int64
15	yr_renovated	21613 non-null	int64
16	zipcode	21613 non-null	int64
17	lat	21613 non-null	float64
18	long	21613 non-null	float64
19	sqft_living15	21613 non-null	int64
20	sqft_lot15	21613 non-null	int64

```
dtypes: float64(5), int64(15), object(1)
```

```
memory usage: 3.5+ MB
```

In [22]:

1 df.describe()

Out[22]:

	id	SalePrice	bedrooms	bathrooms	sqft_living	sqft_lot	floors
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+04	21613.000000
mean	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+04	1.494309
std	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+04	0.539989
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02	1.000000
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000

In [23]:

1 df.head()

Out[23]:

	id	date	SalePrice	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0	0
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0

5 rows × 21 columns

In [24]:


1 pd.isnull(df).any()

Out[24]:

id	False
date	False
SalePrice	False
bedrooms	False
bathrooms	False
sqft_living	False
sqft_lot	False
floors	False
waterfront	False
view	False
condition	False
grade	False
sqft_above	False
sqft_basement	False
yr_built	False
yr_renovated	False
zipcode	False
lat	False
long	False
sqft_living15	False
sqft_lot15	False
dtype:	bool


```
In [25]: 1 plt.style.use('seaborn-whitegrid')
2 fig = plt.figure(figsize=(20,4))
3 sns.countplot(y="grade", data=df)
```

Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x258a45c1208>




```
In [26]: 1 plt.style.use('seaborn-whitegrid')
2 fig = plt.figure(figsize=(20,4))
3 sns.countplot(y="yr_built", data=df)
```

Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x258a45db048>



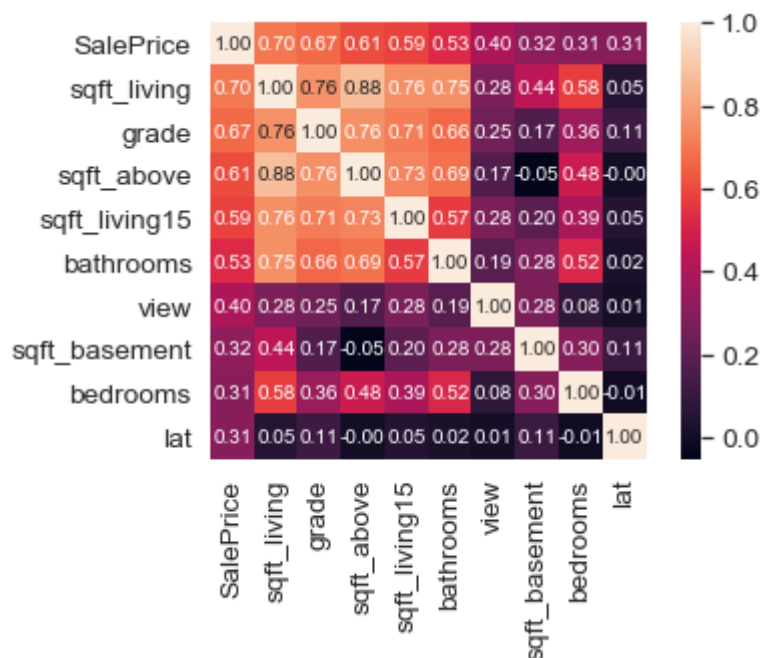
In [27]:

```
1 #correlation matrix
2 corrmat = df.corr()
3 f, ax = plt.subplots(figsize=(12, 9))
4 sns.heatmap(corrmat, vmax=.8, square=True);
```



In [28]:

```
1 #saleprice corr
2 k = 10
3 cols = corrmat.nlargest(k, 'SalePrice')['SalePrice'].index
4 cm = np.corrcoef(df[cols].values.T)
5 sns.set(font_scale=1.25)
6 hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'siz
7 plt.show()
```



In [5]:

```
1 y = df['SalePrice']
2 X = df.drop(['SalePrice', 'date'], axis = 1)
3 names = list(X.columns.values)
4 names
```

Out[5]:

```
['id',
 'bedrooms',
 'bathrooms',
 'sqft_living',
 'sqft_lot',
 'floors',
 'waterfront',
 'view',
 'condition',
 'grade',
 'sqft_above',
 'sqft_basement',
 'yr_built',
 'yr_renovated',
 'zipcode',
 'lat',
 'long',
 'sqft_living15',
 'sqft_lot15']
```

```
In [6]: 1 #introducing missing values
2 for i in range((int)(X.size * 0.1)):
3     row_index = np.random.randint(X.shape[0])
4     col_index = np.random.randint(X.shape[1])
5     col = X.columns[col_index]
6     X.loc[row_index,col] = np.nan
7 # Check what percentage of the data is missing
8 val = 0
9 for col in X.columns:
10     val += X[col].count()
11
12 print(val / X.size)

0.9048331048321308
```

```
In [7]: 1 X.isnull().sum()
```

```
Out[7]: id                1940
bedrooms            2098
bathrooms           2017
sqft_living         2072
sqft_lot            2010
floors              1971
waterfront          2146
view                2054
condition            1993
grade               2057
sqft_above          2083
sqft_basement       2064
yr_built            2125
yr_renovated        2052
zipcode             2096
lat                 2077
long                2061
sqft_living15       2102
sqft_lot15          2062
dtype: int64
```

```
In [45]: 1 X.head()
```

Out[45]:

	id	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_living15
0	7.129301e+09	3.0	NaN	1180.0	5650.0	1.0	NaN	NaN	3.0	7.0	1113.6
1	6.414100e+09	3.0	2.25	2570.0	7242.0	2.0	0.0	0.0	3.0	7.0	2570.0
2	5.631500e+09	2.0	1.00	770.0	10000.0	1.0	0.0	0.0	3.0	6.0	770.0
3	2.487201e+09	4.0	NaN	1960.0	5000.0	1.0	0.0	0.0	5.0	NaN	1960.0
4	1.954401e+09	3.0	2.00	1680.0	NaN	1.0	0.0	0.0	3.0	8.0	1680.0

In [8]:

```
1 #Imputing the missing values with median
2 X = X.apply(lambda x: x.fillna(x.mean()),axis=0)
3 X.head()
```

Out[8]:

	id	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	gr
0	7.129301e+09	3.000000	1.00	1180.0	5650.000000	1.492058	0.000000	0.0	3.0	
1	6.414100e+09	3.000000	2.25	2570.0	7242.000000	2.000000	0.007449	0.0	3.0	
2	4.578891e+09	3.369408	1.00	770.0	10000.000000	1.000000	0.000000	0.0	3.0	
3	2.487201e+09	4.000000	3.00	1960.0	15133.180483	1.000000	0.000000	0.0	5.0	
4	1.954401e+09	3.000000	2.00	1680.0	8080.000000	1.000000	0.000000	0.0	3.0	

In [9]:

```
1 total = X.isnull().sum().sort_values(ascending=False)
2 percent = (X.isnull().sum()/X.isnull().count()).sort_values(ascending=False)
3 missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
4 missing_data.head(20)
```

Out[9]:

	Total	Percent
sqft_lot15	0	0.0
condition	0	0.0
bedrooms	0	0.0
bathrooms	0	0.0
sqft_living	0	0.0
sqft_lot	0	0.0
floors	0	0.0
waterfront	0	0.0
view	0	0.0
grade	0	0.0
sqft_living15	0	0.0
sqft_above	0	0.0
sqft_basement	0	0.0
yr_built	0	0.0
yr_renovated	0	0.0
zipcode	0	0.0
lat	0	0.0
long	0	0.0
id	0	0.0

Scaling


```
In [10]: 1 from sklearn.preprocessing import MinMaxScaler
2 from sklearn.model_selection import train_test_split
3
4 X_train_org, X_test_org, y_train, y_test = train_test_split(X,y, random_state = 0)
5
6 scaler = MinMaxScaler()
7 X_train = scaler.fit_transform(X_train_org)
8 X_test = scaler.transform(X_test_org)
```

```
In [11]: 1 import warnings
2 warnings.filterwarnings("ignore", category=FutureWarning)
3 warnings.filterwarnings("ignore", category=DeprecationWarning)
4 from sklearn.exceptions import ConvergenceWarning
5 from warnings import filterwarnings
6 filterwarnings('ignore')
```

Bagging

Decision Tree Regressor

```
In [54]: 1 from sklearn.ensemble import BaggingRegressor
2 from sklearn.tree import DecisionTreeRegressor
3 from sklearn.model_selection import GridSearchCV
4
5 mf = [2, 5, 10]
6 n = [100, 200, 300, 500]
7 ms = [0.1, 0.5, 1]
8
9 param_grid = dict(max_features = mf, n_estimators = n, max_samples = ms)
10
11 dt_rgrsr = DecisionTreeRegressor(max_depth = 5, random_state=0)
12 bag_rgrsr = GridSearchCV(BaggingRegressor(dt_rgrsr, bootstrap=True, oob_score=True),
13 param_grid)
14 bag_rgrsr.fit(X_train, y_train)
15 y_pred = bag_rgrsr.predict(X_test)
16 print("Best Hyper Parameters:", bag_rgrsr.best_params_)
```

Best Hyper Parameters: {'max_features': 10, 'max_samples': 0.5, 'n_estimators': 200}

```
In [56]: 1 print('Train score: {:.4f}'.format(bag_rgrsr.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(bag_rgrsr.score(X_test, y_test)))
```

Train score: 0.7491

Test score: 0.7212

Linear Regressor

```
In [57]: 1 from sklearn.linear_model import LinearRegression
2
3 lregrsr = LinearRegression()
4
5 bag_rgrsr = GridSearchCV(BaggingRegressor(lregrsr, bootstrap=True, oob_score=True),
6
7 bag_rgrsr.fit(X_train, y_train)
8 y_pred = bag_rgrsr.predict(X_test)
9 bag_rgrsr.best_params_
10 print("Best Hyper Parameters:", bag_rgrsr.best_params_)
```

Best Hyper Parameters: {'max_features': 10, 'max_samples': 0.5, 'n_estimators': 200}

```
In [58]: 1 print('Train score: {:.4f}'.format(bag_rgrsr.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(bag_rgrsr.score(X_test, y_test)))
```

Train score: 0.6482

Test score: 0.6280

Pasting

Decision Tree Regression

```
In [59]: 1 from sklearn.ensemble import BaggingRegressor
2 from sklearn.tree import DecisionTreeRegressor
3 from sklearn.model_selection import GridSearchCV
4
5 mf = [2, 5, 10]
6 n = [100, 200, 300, 500]
7 ms = [0.1, 0.5, 1]
8
9 param_grid = dict(max_features = mf, n_estimators = n, max_samples = ms)
10
11 dt_rgrsr = DecisionTreeRegressor(max_depth = 5, random_state=0)
12 bag_rgrsr = GridSearchCV(BaggingRegressor(dt_rgrsr, bootstrap=False, random_state=0),
13
14 bag_rgrsr.fit(X_train, y_train)
15 y_pred = bag_rgrsr.predict(X_test)
16 print("Best Hyper Parameters:", bag_rgrsr.best_params_)
```

Best Hyper Parameters: {'max_features': 10, 'max_samples': 0.5, 'n_estimators': 200}

```
In [62]: 1 print('Train score: {:.4f}'.format(bag_rgrsr.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(bag_rgrsr.score(X_test, y_test)))
```

Train score: 0.7551

Test score: 0.7232

Linear Regression

```
In [63]: 1 from sklearn.linear_model import LinearRegression
2
3 lreg = LinearRegression()
4
5 bag_rgrsr = GridSearchCV(BaggingRegressor(lreg, bootstrap=False, random_state=0),
6
7 bag_rgrsr.fit(X_train, y_train)
8 y_pred = bag_rgrsr.predict(X_test)
9 print("Best Hyper Parameters:", bag_rgrsr.best_params_)
```

Best Hyper Parameters: {'max_features': 10, 'max_samples': 0.1, 'n_estimators': 200}

```
In [64]: 1 print('Train score: {:.4f}'.format(bag_rgrsr.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(bag_rgrsr.score(X_test, y_test)))
```

Train score: 0.6491

Test score: 0.6287

Adaboost Boosting

Decision Tree Regressor

```
In [65]: 1 from sklearn.ensemble import AdaBoostRegressor
2
3
4 n = [100, 200, 300, 500]
5 lr = [0.1, 0.5, 1]
6
7 param_grid = dict(n_estimators = n, learning_rate = lr)
8
9 dt_rgrsr = DecisionTreeRegressor(max_depth = 5, random_state=0)
10 ada_rgrsr = GridSearchCV(AdaBoostRegressor(dt_rgrsr, random_state=0), param_grid, c
11 ada_rgrsr.fit(X_train, y_train)
12 y_pred = ada_rgrsr.predict(X_test)
13 print("Best Hyper Parameters:", ada_rgrsr.best_params_)
```

Best Hyper Parameters: {'learning_rate': 0.1, 'n_estimators': 100}

```
In [66]: 1 print('Train score: {:.4f}'.format(ada_rgrsr.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(ada_rgrsr.score(X_test, y_test)))
```

Train score: 0.8016

Test score: 0.7597

Linear regression

```
In [67]: 1 n = [100, 200, 300, 500]
2 lr = [0.1, 0.5, 1]
3
4 param_grid = dict(n_estimators = n, learning_rate = lr)
5
6 ada_rgrsr = GridSearchCV(AdaBoostRegressor(lreg, random_state=0), param_grid, cv =
7 ada_rgrsr.fit(X_train, y_train)
8 y_pred = ada_rgrsr.predict(X_test)
9 print("Best Hyper Parameters:", ada_rgrsr.best_params_)
```

Best Hyper Parameters: {'learning_rate': 0.1, 'n_estimators': 100}

```
In [68]: 1 print('Train score: {:.4f}'.format(ada_rgrsr.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(ada_rgrsr.score(X_test, y_test)))
```

Train score: 0.4398

Test score: 0.4299

Gradient Boosting

```
In [69]: 1 from sklearn.ensemble import GradientBoostingRegressor
2
3 mf = [2, 5, 10]
4 n = [50, 100, 200]
5 md = [1, 5, 10]
6
7 param_grid = dict(max_features = mf, n_estimators = n, max_depth = md)
8
9 grdbst = GridSearchCV(GradientBoostingRegressor(random_state=0), param_grid, cv = 5
10
11 grdbst.fit(X_train, y_train)
12 y_pred = grdbst.predict(X_test)
13 print("Best Hyper Parameters:", grdbst.best_params_)
```

Best Hyper Parameters: {'max_depth': 5, 'max_features': 5, 'n_estimators': 200}

```
In [70]: 1 print('Train score: {:.4f}'.format(grdbst.score(X_train, y_train)))
2 print('Test score: {:.4f}'.format(grdbst.score(X_test, y_test)))
```

Train score: 0.9404

Test score: 0.8654

Principal Component Analysis

```
In [71]: 1 #reducing dimensions using PCA to create new dataset
2 from sklearn.decomposition import PCA
3
4 pca = PCA(n_components = 0.95)
5 X_train_reduced = pca.fit_transform(X_train)
6 X_test_reduced = pca.transform(X_test)
7
8 pca.n_components_
```

Out[71]: 11

Linear Regression without PCA

```
In [79]: 1 lregrsr = LinearRegression()
2 lregrsr.fit(X_train, y_train)
3 print(lregrsr.score(X_train, y_train))
4 print(lregrsr.score(X_test, y_test))
```

```
0.674242240793372
0.6557793731457227
```

Linear Regression with PCA

```
In [80]: 1 lregrsr.fit(X_train_reduced, y_train)
2 print(lregrsr.score(X_train_reduced, y_train))
3 print(lregrsr.score(X_test_reduced, y_test))
```

```
0.6443319181355665
0.6155325953853179
```

PCA reduced test and train scores for linear regression

Polynomial Regression without PCA

```
In [75]: 1 from sklearn.preprocessing import PolynomialFeatures
2
3 train_score = []
4 test_score = []
5
6 for n in range(1,3):
7     polyfts = PolynomialFeatures(n)
8     X_train_poly = polyfts.fit_transform(X_train)
9     X_test_poly = polyfts.transform(X_test)
10    lregrsr.fit(X_train_poly, y_train)
11    train_score_list.append(lregrsr.score(X_train_poly, y_train))
12    test_score_list.append(lregrsr.score(X_test_poly, y_test))
```

```
In [76]: 1 print(train_score)
2 print(test_score)
```

```
[0.674242240793372, 0.7966594924663959]
[0.6557793731457247, 0.7747840096656062]
```

Polynomial Regression without PCA

```
In [77]: 1 train_score = []
2 test_score = []
3
4 for n in range(1,3):
5     poly = PolynomialFeatures(n)
6     X_train_poly = poly.fit_transform(X_train_reduced)
7     X_test_poly = poly.transform(X_test_reduced)
8     lregrsr.fit(X_train_poly, y_train)
9     train_score_list.append(lregrsr.score(X_train_poly, y_train))
10    test_score_list.append(lregrsr.score(X_test_poly, y_test))
```

```
In [78]: 1 print(train_score)
2 print(test_score)
```

```
[0.6443319181355665, 0.7553138111248706]
[0.6155325953853179, 0.7193217271924304]
```

Implementing PCA reduced the testing and training scores for Polynomial regression

Ridge without PCA

```
In [83]: 1 from sklearn.linear_model import Ridge
2
3 alpha = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
4 param_grid = dict(alpha=alpha)
5
6 ridge_regrsr = GridSearchCV(Ridge(), param_grid=param_grid, scoring='r2', verbose=1)
7 ridge_regrsr.fit(X_train, y_train)
8
9 y_pred = ridge_regrsr.predict(X_test)
10
11 print("Best Hyper Parameters:",ridge_regrsr.best_params_)
```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

Best Hyper Parameters: {'alpha': 1}

[Parallel(n_jobs=-1)]: Done 35 out of 35 | elapsed: 2.5s finished

```
In [84]: 1 print('Train score: {:.4f}'.format(ridge_regrsr.score(X_train,y_train)))
2 print('Test score: {:.4f}'.format(ridge_regrsr.score(X_test, y_test)))
```

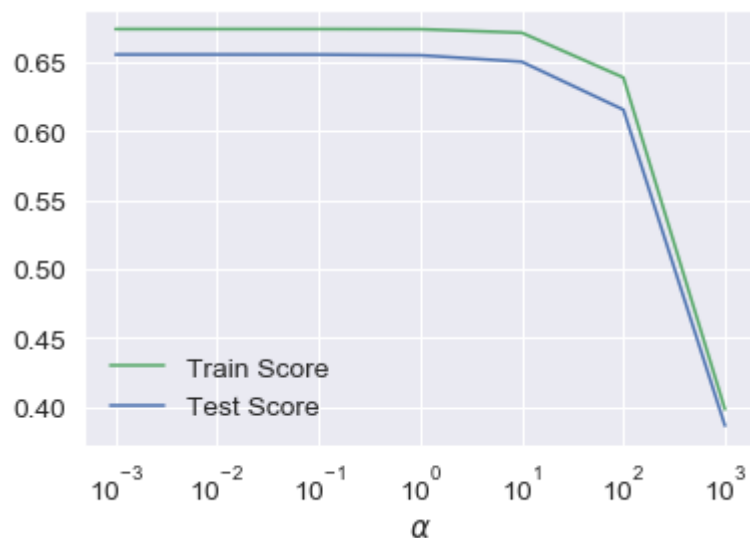
```
Train score: 0.6742
Test score: 0.6552
```

```

In [85]: 1 #variation of scores with alpha
2 train_score_list = []
3 test_score_list = []
4
5 x_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
6 for alpha in x_range:
7     ridge = Ridge(alpha)
8     ridge.fit(X_train,y_train)
9     train_score_list.append(ridge.score(X_train,y_train))
10    test_score_list.append(ridge.score(X_test, y_test))
11
12 %matplotlib inline
13 import matplotlib.pyplot as plt
14 plt.plot(x_range, train_score_list, c = 'g', label = 'Train Score')
15 plt.plot(x_range, test_score_list, c = 'b', label = 'Test Score')
16 plt.xscale('log')
17 plt.legend(loc = 3)
18 plt.xlabel(r'$\alpha$')

```

Out[85]: Text(0.5, 0, '\$\alpha\$')



Ridge with PCA

```
In [87]: 1 ridge_regrsr.fit(X_train_reduced,y_train)
2
3 y_pred = ridge_regrsr.predict(X_test_reduced)
4
5 print("Best Hyper Parameters:",ridge_regrsr.best_params_)
```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

Best Hyper Parameters: {'alpha': 1}

[Parallel(n_jobs=-1)]: Done 20 out of 35 | elapsed: 0.1s remaining: 0.1s

[Parallel(n_jobs=-1)]: Done 35 out of 35 | elapsed: 0.2s finished

```
In [88]: 1 print('Train score: {:.4f}'.format(ridge_regrsr.score(X_train_reduced,y_train)))
2 print('Test score: {:.4f}'.format(ridge_regrsr.score(X_test_reduced, y_test)))
```

Train score: 0.6443

Test score: 0.6156

```
In [89]: 1 #variation scores with alpha
2 train_score_list = []
3 test_score_list = []
4
5 x_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
6 for alpha in x_range:
7     ridge = Ridge(alpha)
8     ridge.fit(X_train_reduced,y_train)
9     train_score_list.append(ridge.score(X_train_reduced,y_train))
10    test_score_list.append(ridge.score(X_test_reduced, y_test))
11
12 %matplotlib inline
13 import matplotlib.pyplot as plt
14 plt.plot(x_range, train_score_list, c = 'g', label = 'Train Score')
15 plt.plot(x_range, test_score_list, c = 'b', label = 'Test Score')
16 plt.xscale('log')
17 plt.legend(loc = 3)
18 plt.xlabel(r'$\alpha$')
```

Out[89]: Text(0.5, 0, '\$\alpha\$')



With PCA the best parameters for ridge varied and also the scores got reduced.

Lasso without PCA

In [90]:

```
1 from sklearn.linear_model import Lasso
2
3 lasso = Lasso(alpha = 0.01)
4 lasso.fit(X_train,y_train)
5
6 alpha = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
7 param_grid = dict(alpha=alpha)
8
9 lasso_regrsr = GridSearchCV(estimator=lasso, param_grid=param_grid, scoring='r2', v
10 lasso_regrsr.fit(X_train, y_train)
11
12 y_pred = lasso_regrsr.predict(X_test)
13 print("Best Hyper Parameters:",lasso_regrsr.best_params_)
14
```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

Best Hyper Parameters: {'alpha': 1}

[Parallel(n_jobs=-1)]: Done 20 out of 35 | elapsed: 0.4s remaining: 0.3s

[Parallel(n_jobs=-1)]: Done 35 out of 35 | elapsed: 0.5s finished

In [91]:

```
1 print('Train score: {:.4f}'.format(lasso_regrsr.score(X_train,y_train)))
2 print('Test score: {:.4f}'.format(lasso_regrsr.score(X_test, y_test)))
```

Train score: 0.6742

Test score: 0.6558

```

In [92]: 1 train_score_list = []
          2 test_score_list = []
          3
          4 x_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
          5 for alpha in x_range:
          6     lasso = Lasso(alpha)
          7     lasso.fit(X_train,y_train)
          8     train_score_list.append(lasso.score(X_train,y_train))
          9     test_score_list.append(lasso.score(X_test, y_test))
         10
         11 plt.plot(x_range, train_score_list, c = 'g', label = 'Train Score')
         12 plt.plot(x_range, test_score_list, c = 'b', label = 'Test Score')
         13 plt.xscale('log')
         14 plt.legend(loc = 3)
         15 plt.xlabel(r'$\alpha$')

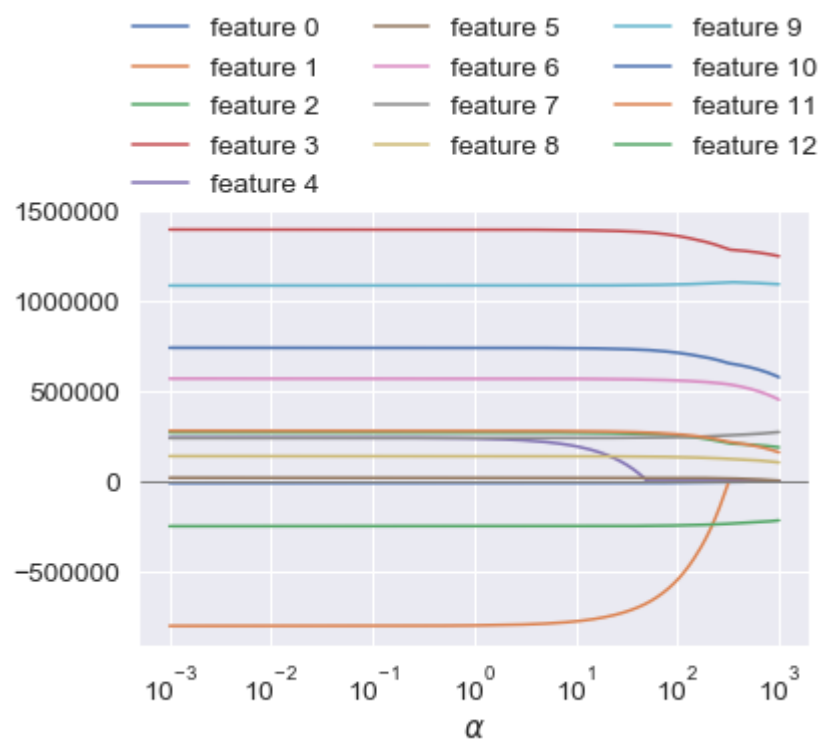
```

Out[92]: Text(0.5, 0, '\$\alpha\$')



In [96]:

```
1 %matplotlib inline
2
3 x_range1 = np.linspace(0.001, 1, 1000).reshape(-1,1)
4 x_range2 = np.linspace(1, 1000, 1000).reshape(-1,1)
5
6 x_range = np.append(x_range1, x_range2)
7 coeff = []
8
9 for alpha in x_range:
10     lasso = Lasso(alpha)
11     lasso.fit(X_train,y_train)
12     coeff.append(lasso.coef_ )
13
14 coeff = np.array(coeff)
15
16 for i in range(0,13):
17     plt.plot(x_range, coeff[:,i], label = 'feature {:d}'.format(i))
18
19 plt.axhline(y=0, xmin=0.001, xmax=9999, linewidth=1, c = 'gray')
20 plt.xlabel(r'$\alpha$')
21 plt.xscale('log')
22 plt.legend(loc='upper center', bbox_to_anchor=(0.5, 1.5),
23           ncol=3, fancybox=True, shadow=True)
24 plt.show()
```



Lasso with PCA

```
In [98]: 1 lasso_regrsr.fit(X_train_reduced, y_train)
2
3 y_pred = lasso_regrsr.predict(X_test_reduced)
4
5 print("Best Hyper Parameters:",lasso_regrsr.best_params_)
```

Fitting 5 folds for each of 7 candidates, totalling 35 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

Best Hyper Parameters: {'alpha': 100}

[Parallel(n_jobs=-1)]: Done 35 out of 35 | elapsed: 1.9s finished

```
In [99]: 1 print('Train score: {:.4f}'.format(lasso_regrsr.score(X_train_reduced,y_train)))
2 print('Test score: {:.4f}'.format(lasso_regrsr.score(X_test_reduced, y_test)))
```

Train score: 0.6443

Test score: 0.6155

```
In [100]: 1 train_score_list = []
2 test_score_list = []
3
4 x_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
5 for alpha in x_range:
6     lasso = Lasso(alpha)
7     lasso.fit(X_train_reduced,y_train)
8     train_score_list.append(lasso.score(X_train_reduced,y_train))
9     test_score_list.append(lasso.score(X_test_reduced, y_test))
10
11 plt.plot(x_range, train_score_list, c = 'g', label = 'Train Score')
12 plt.plot(x_range, test_score_list, c = 'b', label = 'Test Score')
13 plt.xscale('log')
14 plt.legend(loc = 3)
15 plt.xlabel(r'$\alpha$')
```

Out[100]: Text(0.5, 0, '\$\alpha\$')



PCA reduced the scores and changed best parameters for Lasso

KNN Regressor without PCA

In [101]:

```
1 from sklearn.neighbors import KNeighborsRegressor
2
3 n = [3,5,11,19]
4
5 param_grid = dict(n_neighbors=n)
6
7 knnr=GridSearchCV(KNeighborsRegressor(), param_grid, cv = 5, iid = False)
8 knnr.fit(X_train,y_train)
9 y_pred = knnr.predict(X_test)
10
11 print("Best Hyper Parameters:",knnr.best_params_)
```

Best Hyper Parameters: {'n_neighbors': 5}

In [102]:

```
1 print('Train score: {:.4f}'.format(knnr.score(X_train,y_train)))
2 print('Test score: {:.4f}'.format(knnr.score(X_test, y_test)))
```

Train score: 0.8062

Test score: 0.7035

KNN Regressor with PCA

In [103]:

```
1 knnr.fit(X_train_reduced,y_train)
2
3 y_pred = knnr.predict(X_test_reduced)
4
5 print("Best Hyper Parameters:",knnr.best_params_)
```

Best Hyper Parameters: {'n_neighbors': 5}

In [104]:

```
1 print('Train score: {:.4f}'.format(knnr.score(X_train_reduced,y_train)))
2 print('Test score: {:.4f}'.format(knnr.score(X_test_reduced, y_test)))
```

Train score: 0.7990

Test score: 0.6942

KNN regressor being the best performing model, gave reduced scores after implementing PCA

SVM Regressor without PCA

```
In [105]: 1 from sklearn.svm import SVR
2
3 C = [0.001, 0.01, 0.1, 1, 10, 100]
4 kernel = ['linear', 'poly', 'rbf']
5
6 param_grid = dict(C=C, kernel=kernel)
7
8 SVM=GridSearchCV(SVR(gamma='scale'), param_grid, cv = 5, iid = False)
9
10 SVM.fit(X_train,y_train)
11 y_pred = SVM.predict(X_test)
12
13 print("Best Hyper Parameters:",SVM.best_params_)
```

Best Hyper Parameters: {'C': 100, 'kernel': 'poly'}

```
In [106]: 1 print('Train score: {:.4f}'.format(SVM.score(X_train,y_train)))
2 print('Test score: {:.4f}'.format(SVM.score(X_test, y_test)))
```

Train score: 0.5228
Test score: 0.5079

SVM Regressor with PCA

```
In [108]: 1 SVM.fit(X_train_reduced,y_train)
2
3 y_pred = SVM.predict(X_test_reduced)
4
5 print("Best Hyper Parameters:",SVM.best_params_)
```

Best Hyper Parameters: {'C': 100, 'kernel': 'poly'}

```
In [109]: 1 print('Train score: {:.4f}'.format(SVM.score(X_train_reduced,y_train)))
2 print('Test score: {:.4f}'.format(SVM.score(X_test_reduced, y_test)))
```

Train score: 0.2483
Test score: 0.2184

The best parameters of SVM seems unchanged on implementing PCA but the scores have dropped way down.

Linear SVR without PCA

```
In [110]: 1 from sklearn.svm import SVR, LinearSVR
2
3 C = [0.001, 0.01, 0.1, 1, 10, 100]
4 param_grid = dict(C=C)
5
6 linearsvr=GridSearchCV(SVR(), param_grid, cv = 5, iid = False)
7
8 linearsvr.fit(X_train,y_train)
9 y_pred = linearsvr.predict(X_test)
10
11 print("Best Hyper Parameters:",linearsvr.best_params_)
```

Best Hyper Parameters: {'C': 100}

```
In [111]: 1 print('Train score: {:.4f}'.format(linearsvr.score(X_train,y_train)))
2 print('Test score: {:.4f}'.format(linearsvr.score(X_test, y_test)))
```

Train score: 0.0514

Test score: 0.0614

Linear SVR with PCA

```
In [113]: 1 linearsvr.fit(X_train_reduced,y_train)
2 y_pred = linearsvr.predict(X_test_reduced)
3
4 print("Best Hyper Parameters:",linearsvr.best_params_)
```

Best Hyper Parameters: {'C': 100}

```
In [114]: 1 print('Train score: {:.4f}'.format(linearsvr.score(X_train_reduced,y_train)))
2 print('Test score: {:.4f}'.format(linearsvr.score(X_test_reduced, y_test)))
```

Train score: 0.0485

Test score: 0.0596

GridsearchCV gave similar hyper parameters for Linear SVR after implementing PCA. The scores went down on PCA implementation

Neural Networks

```
In [12]: 1 import keras as k
2 import pandas as pd
```

```
In [13]: 1 X_NNtrain = X_train[:,0:20]
2 X_NNtest = X_test[:,0:20]
```

```
In [14]: 1 X_NNtrain.shape
```

Out[14]: (16209, 19)

```
In [20]: 1 from keras.models import Sequential
2 from keras.layers import Dense
```

```
In [21]: 1 model = Sequential()
2 model.add(Dense(19, input_dim=19, kernel_initializer='normal', activation='relu'))
3 model.add(Dense(1, kernel_initializer='normal'))
```

```
In [22]: 1 model.compile(loss='mse', optimizer='sgd' , metrics = ['mse'])
```

```
In [25]: 1 model.fit(X_NNtrain, y_train, epochs = 100, batch_size = 5)
```

```
Epoch 1/100
 1/3242 [.....] - ETA: 0s - loss: 219795652608.0000 - mse: 219795652608.0000WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0000s vs `on_train_batch_end` time: 0.0010s). Check your callbacks.
3242/3242 [=====] - 2s 654us/step - loss: 256175178971158151168.0000 - mse: 256175178971158151168.0000
Epoch 2/100
3242/3242 [=====] - 2s 642us/step - loss: 135696056320.0000 - mse: 135696056320.0000
Epoch 3/100
3242/3242 [=====] - 2s 691us/step - loss: 135726055424.0000 - mse: 135726055424.0000
Epoch 4/100
3242/3242 [=====] - 2s 709us/step - loss: 135768875008.0000 - mse: 135768875008.0000
Epoch 5/100
3242/3242 [=====] - 2s 698us/step - loss: 135638474752.0000 - mse: 135638491136.0000
Epoch 6/100
```

```
In [24]: 1 model.evaluate(X_NNtest, y_test)
```

```
1/169 [.....] - ETA: 0s - loss: 383676514304.0000 - mse: 383676514304.0000WARNING:tensorflow:Callbacks method `on_test_batch_end` is slow compared to the batch time (batch time: 0.0000s vs `on_test_batch_end` time: 0.0010s). Check your callbacks.
169/169 [=====] - 0s 674us/step - loss: 419171368960.0000 - mse: 419171368960.0000
```

```
Out[24]: [419171368960.0, 419171368960.0]
```

```
In [26]: 1 from sklearn.metrics import r2_score, recall_score, precision_score
2
3 y_train_predict = model.predict(X_NNtrain)
4 y_test_predict = model.predict(X_NNtest)
5
6 print('Train score: {:.2f}'.format(r2_score(y_train, y_train_predict)))
7 print('Test score: {:.2f}'.format(r2_score(y_test, y_test_predict)))
```

```
Train score: -0.00
Test score: -0.00
```